

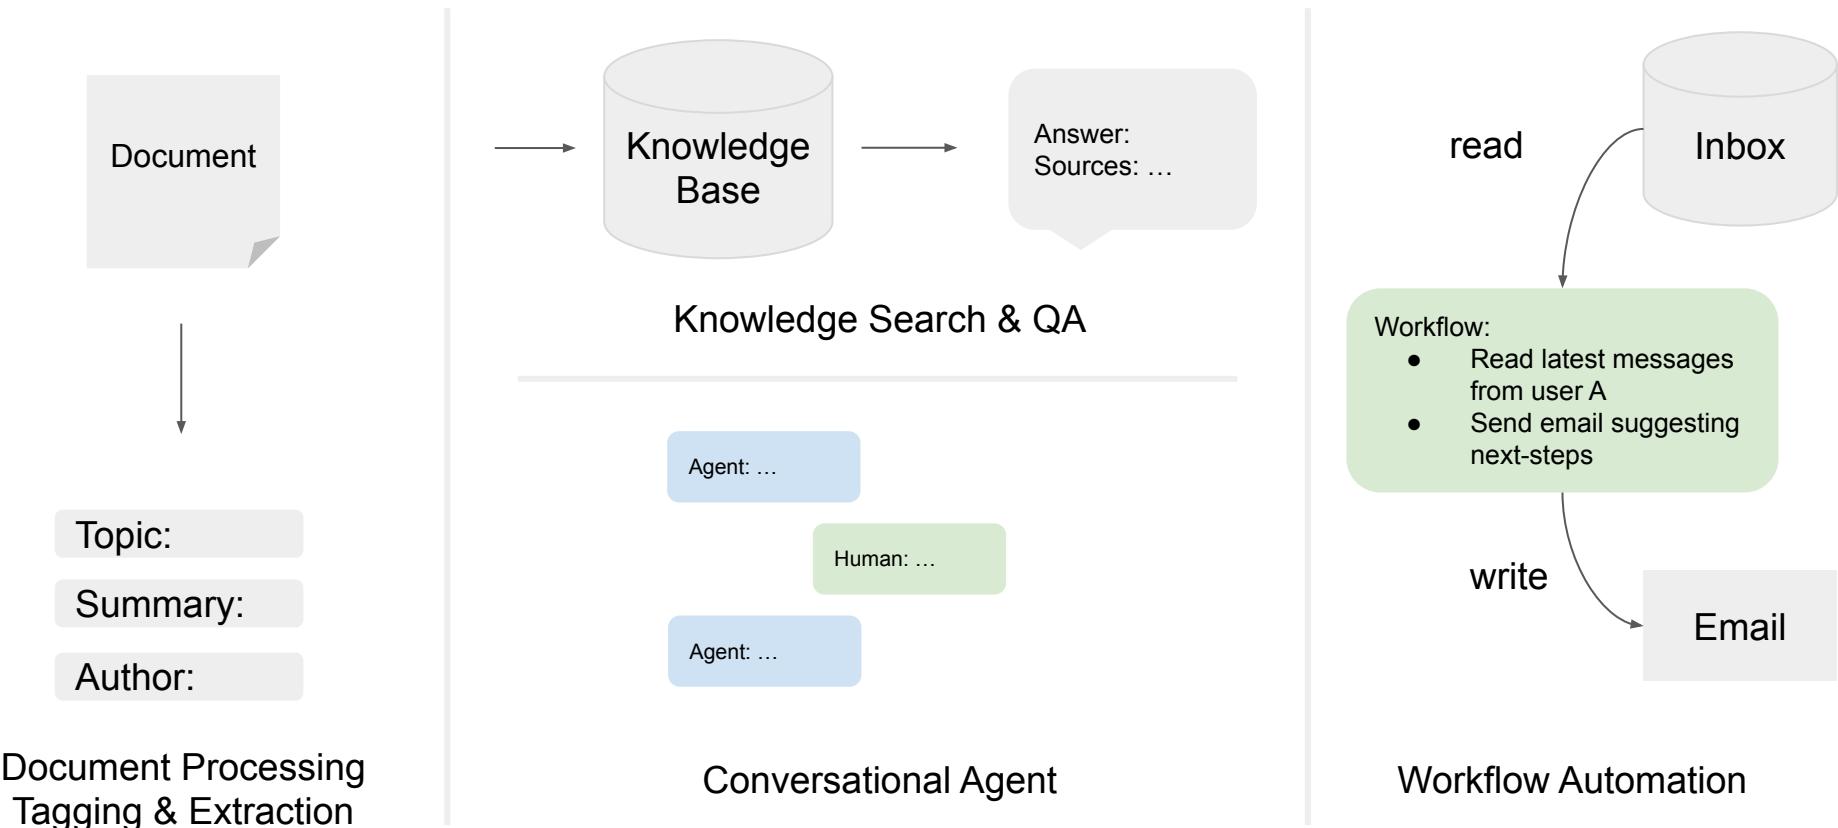


Building and Productionizing RAG

Jerry Liu, LlamaIndex co-founder/CEO



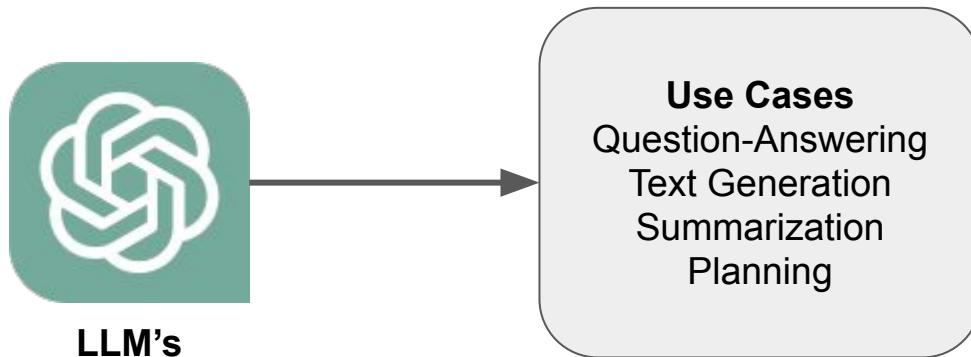
GenAI - Enterprise Use-cases





Context

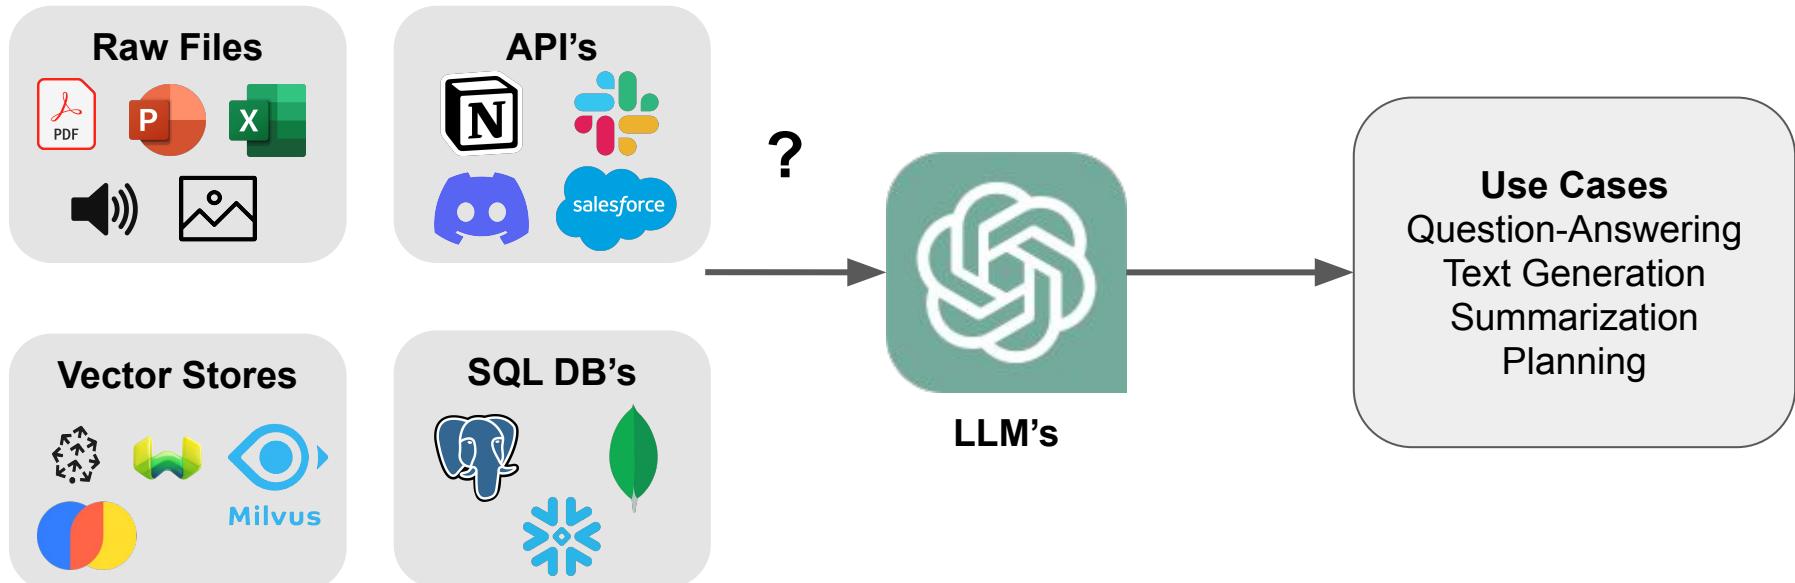
- LLMs are a phenomenal piece of technology for knowledge generation and reasoning. They are pre-trained on large amounts of **publicly available data**.





Context

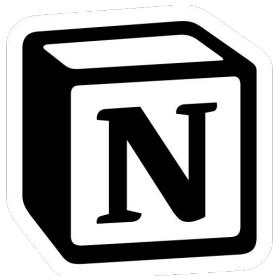
- How do we best augment LLMs with our own **private data**?





Paradigms for inserting knowledge

Retrieval Augmentation - Fix the model, put context into the prompt



Before college the two main things I worked on, outside of school, were writing and programming. I didn't write essays. I wrote what beginning writers were supposed to write then, and probably still are: short stories. My stories were awful. They had hardly any plot, just characters with strong feelings, which I imagined made them deep...

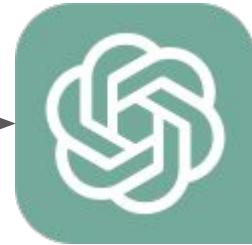


Input Prompt

Here is the context:
Before college the two main things...



Given the context,
answer the following
question:
{query_str}

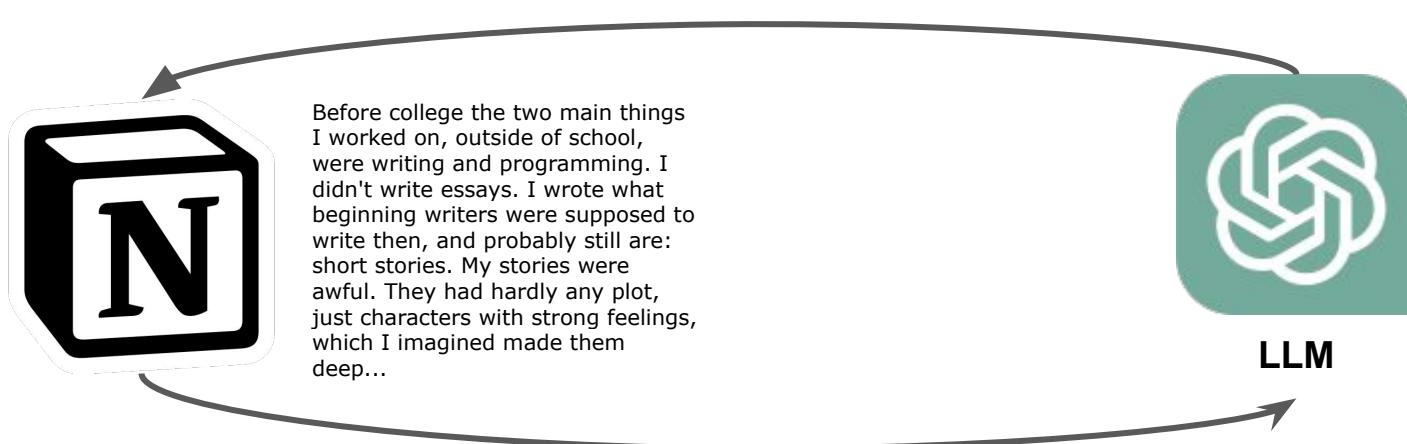


LLM



Paradigms for inserting knowledge

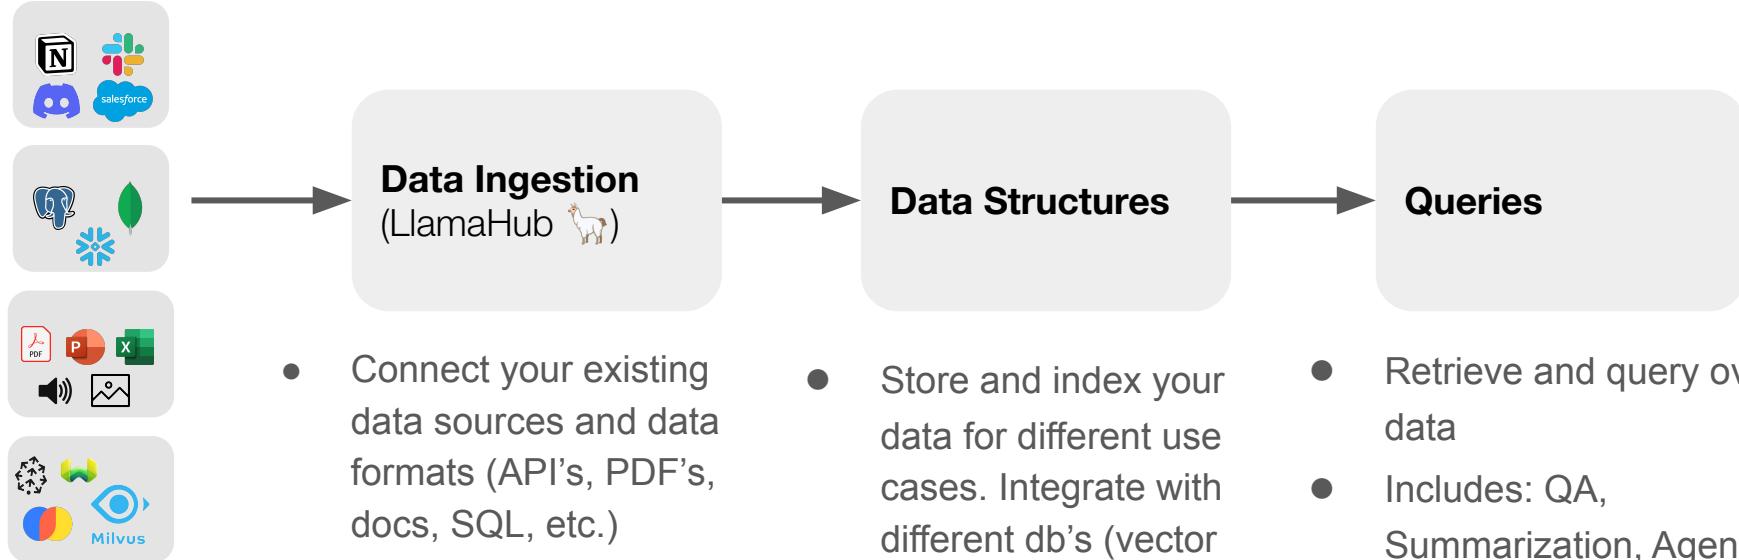
Fine-tuning - baking knowledge into the weights of the network





Llamaindex: A data framework for LLM applications

- Data Management and Query Engine for your LLM application
- Offers components across the data lifecycle: ingest, index, and query over data





quickstart.py

```
from llama_index import VectorStoreIndex, SimpleDirectoryReader

documents = SimpleDirectoryReader('data').load_data()
index = VectorStoreIndex.from_documents(documents)
query_engine = index.as_query_engine()
response = query_engine.query("What did the author do growing up?")
print(str(response))
```



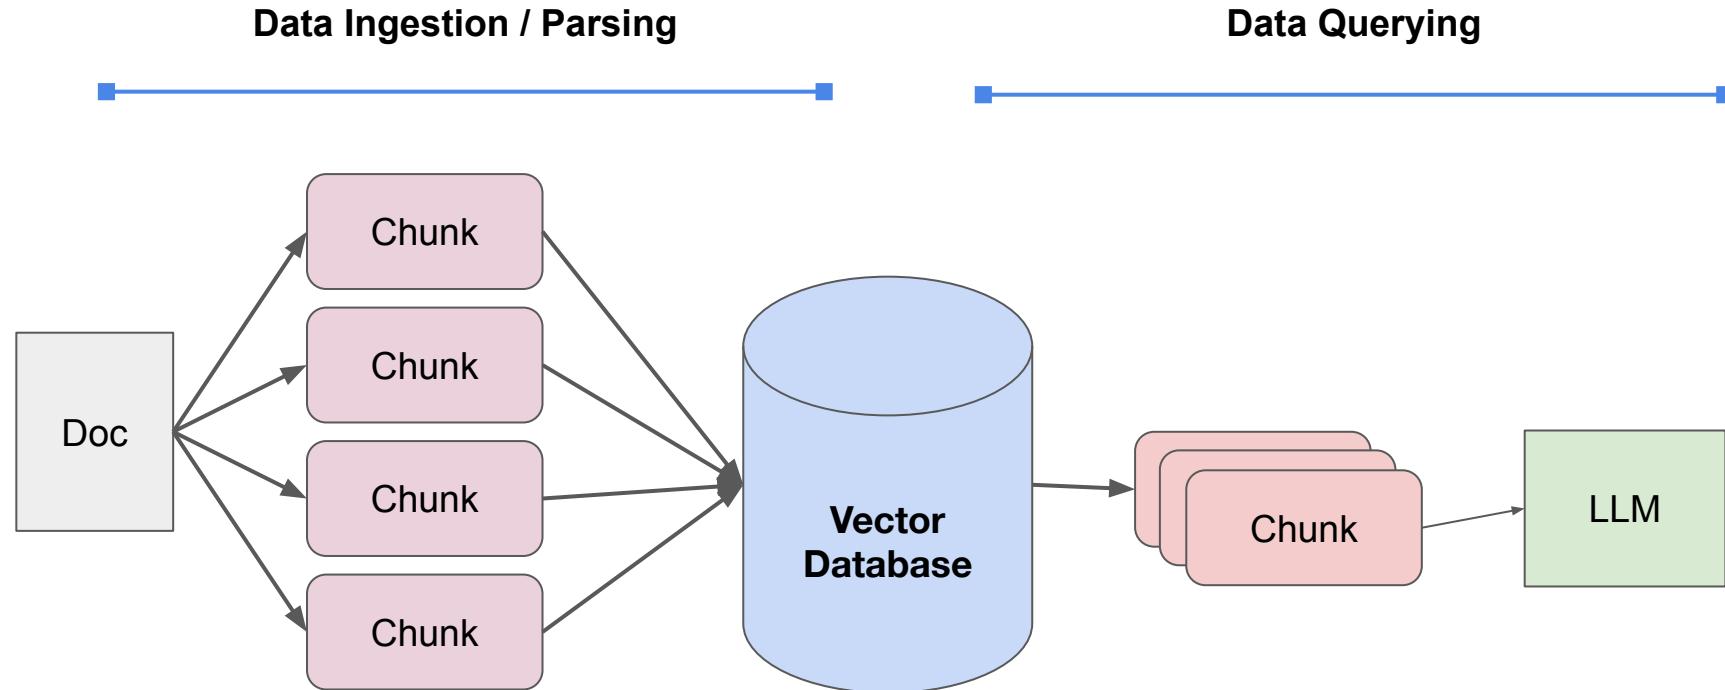
Codelimage



RAG Stack



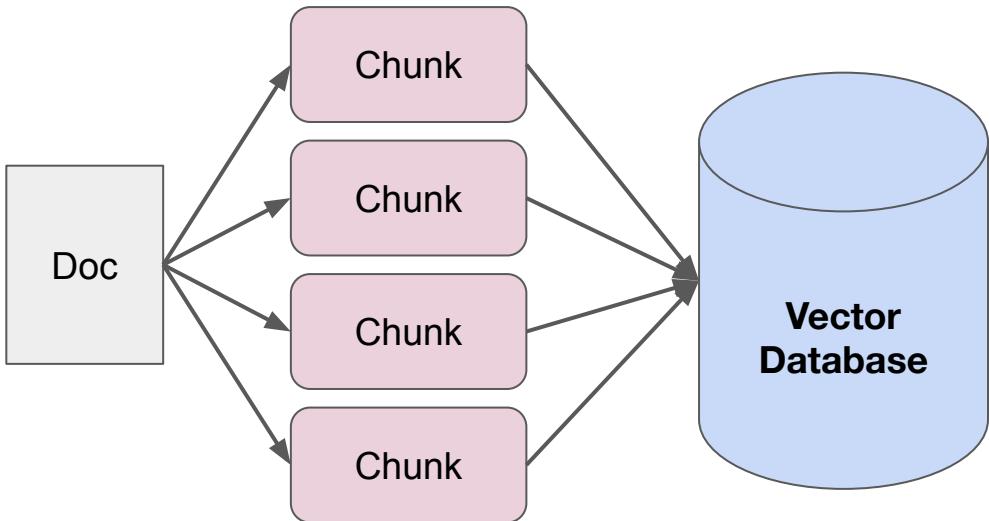
Current RAG Stack for building a QA System



5 Lines of Code in `Llamaindex!`



Current RAG Stack (Data Ingestion/Parsing)



Process:

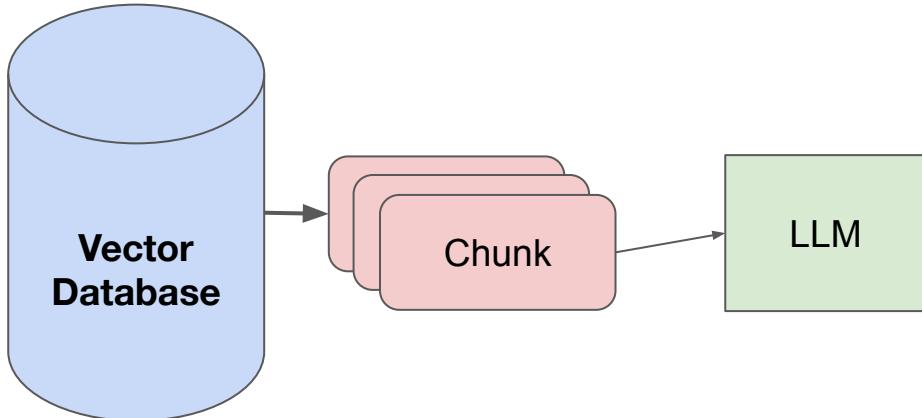
- Split up document(s) into even chunks.
- Each chunk is a piece of raw text.
- Generate embedding for each chunk (e.g. OpenAI embeddings, sentence_transformer)
- Store each chunk into a vector database



Current RAG Stack (Querying)

Process:

- Find top-k most similar chunks from vector database collection
- Plug into LLM response synthesis module

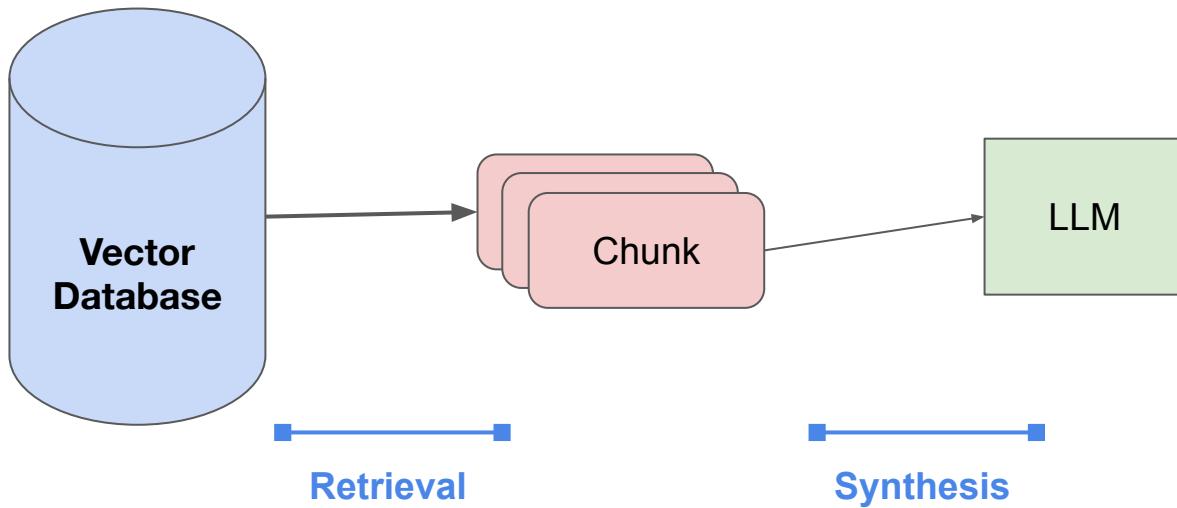




Current RAG Stack (Querying)

Process:

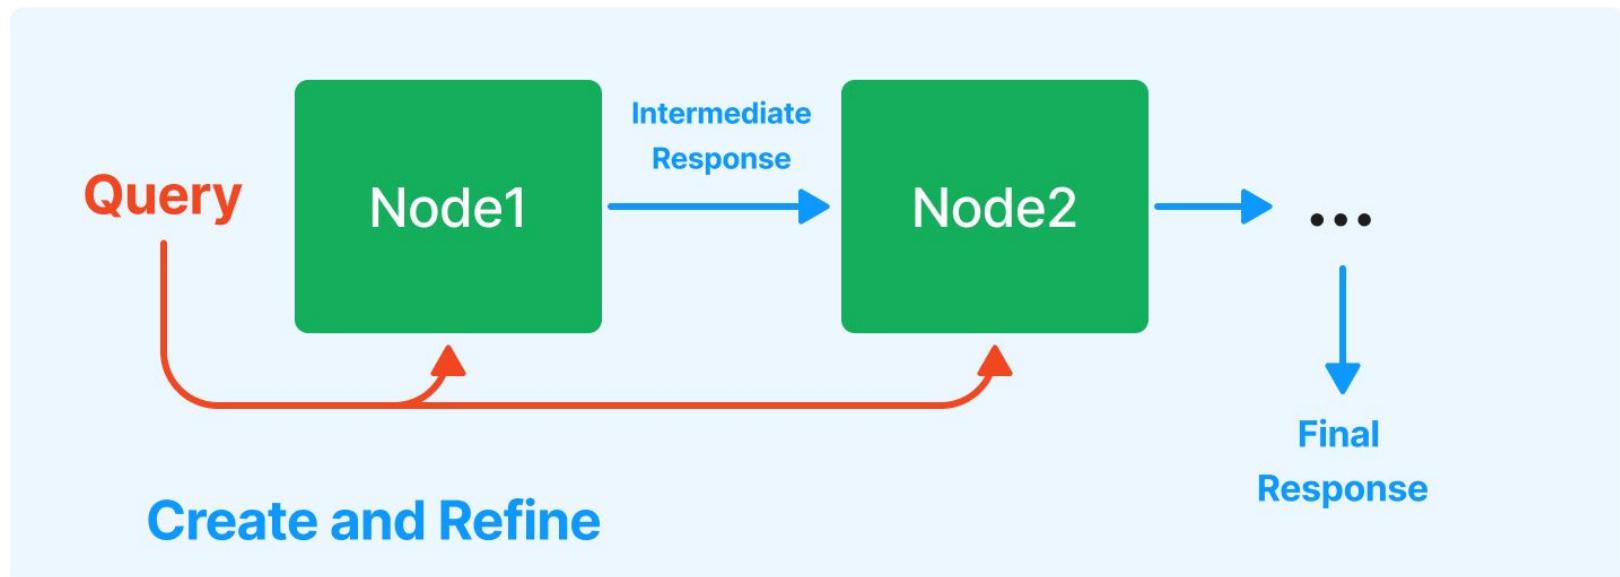
- Find top-k most similar chunks from vector database collection
- Plug into LLM **response synthesis module**





Response Synthesis

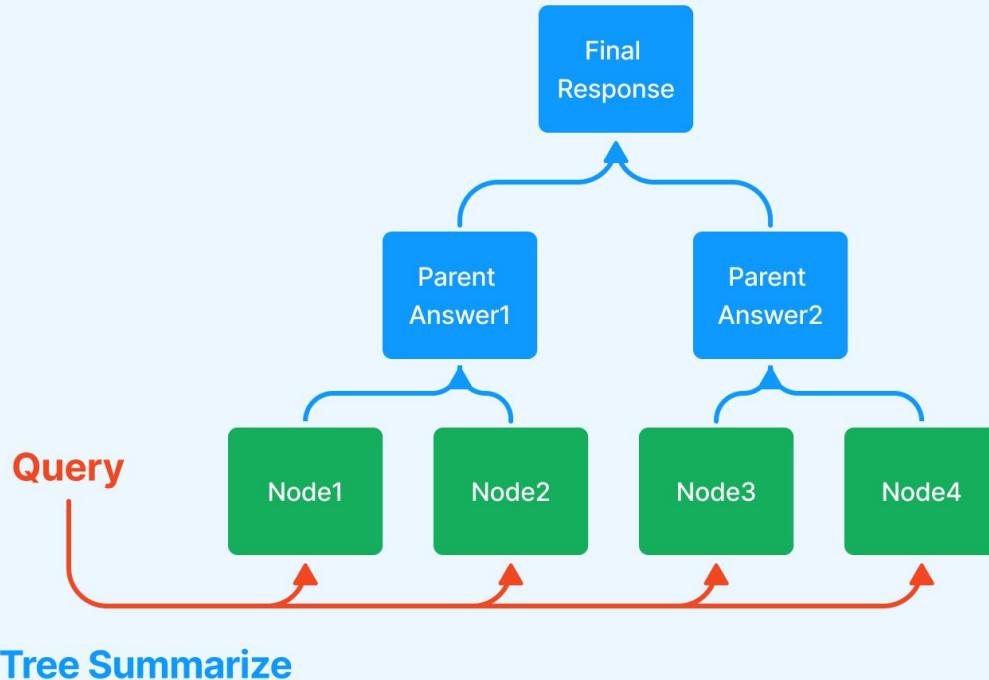
Create and refine





Response Synthesis

Tree Summarize





Quickstart

https://colab.research.google.com/drive/1knQpGJLHj-LTTHqlZhgcjDH5F_nJliY0?usp=sharing





Challenges with “Naive” RAG



Challenges with Naive RAG

- Failure Modes
 - Quality-Related (Hallucination, Accuracy)
 - Non-Quality-Related (Latency, Cost, Syncing)



Challenges with Naive RAG (Response Quality)

- Bad Retrieval
 - **Low Precision:** Not all chunks in retrieved set are relevant
 - Hallucination + Lost in the Middle Problems
 - **Low Recall:** Not all relevant chunks are retrieved.
 - Lacks enough context for LLM to synthesize an answer
 - **Outdated information:** The data is redundant or out of date.



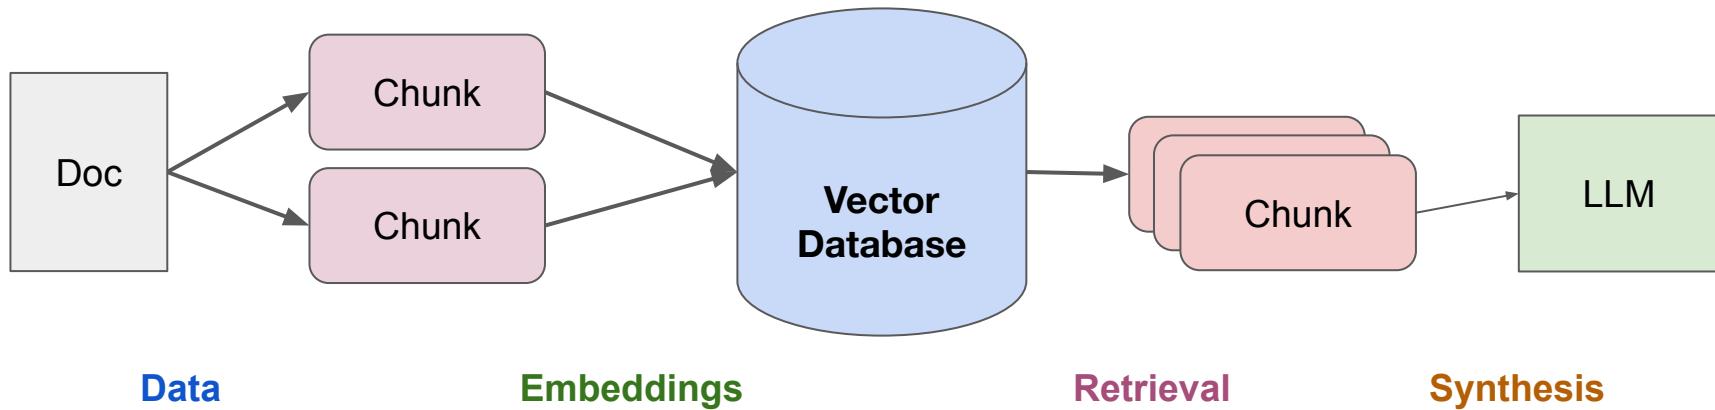
Challenges with Naive RAG (Response Quality)

- Bad Retrieval
 - **Low Precision:** Not all chunks in retrieved set are relevant
 - Hallucination + Lost in the Middle Problems
 - **Low Recall:** Not all relevant chunks are retrieved.
 - Lacks enough context for LLM to synthesize an answer
 - **Outdated information:** The data is redundant or out of date.
- Bad Response Generation
 - **Hallucination:** Model makes up an answer that isn't in the context.
 - **Irrelevance:** Model makes up an answer that doesn't answer the question.
 - **Toxicity/Bias:** Model makes up an answer that's harmful/offensive.



What do we do?

- **Data:** Can we store additional information beyond raw text chunks?
- **Embeddings:** Can we optimize our embedding representations?
- **Retrieval:** Can we do better than top-k embedding lookup?
- **Synthesis:** Can we use LLMs for more than generation?





What do we do?

- **Data:** Can we store additional information beyond raw text chunks?
- **Embeddings:** Can we optimize our embedding representations?
- **Retrieval:** Can we do better than top-k embedding lookup?
- **Synthesis:** Can we use LLMs for more than generation?

But before all this...

We need a way to measure performance

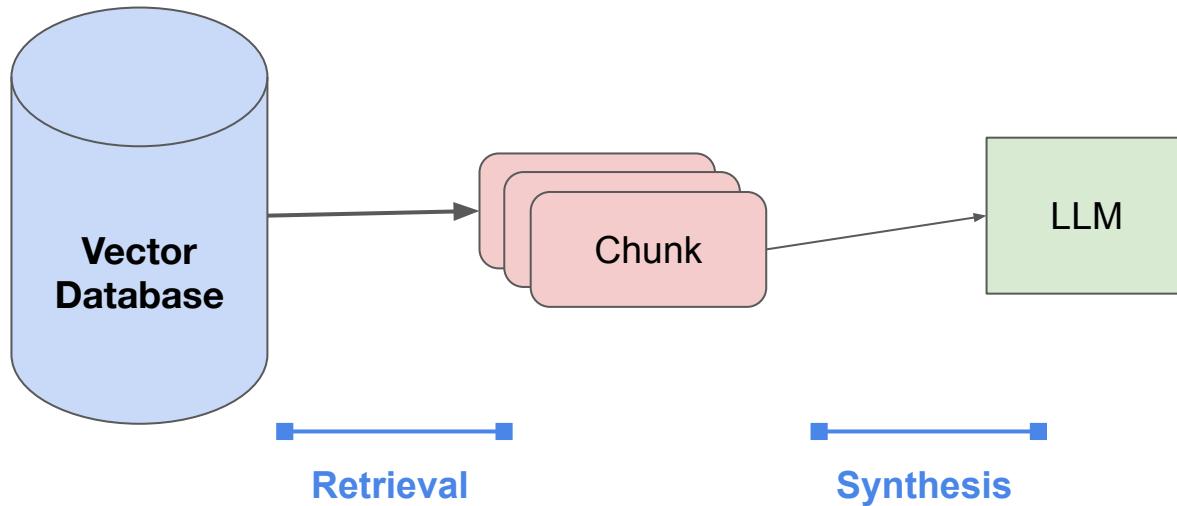


Evaluation



Evaluation

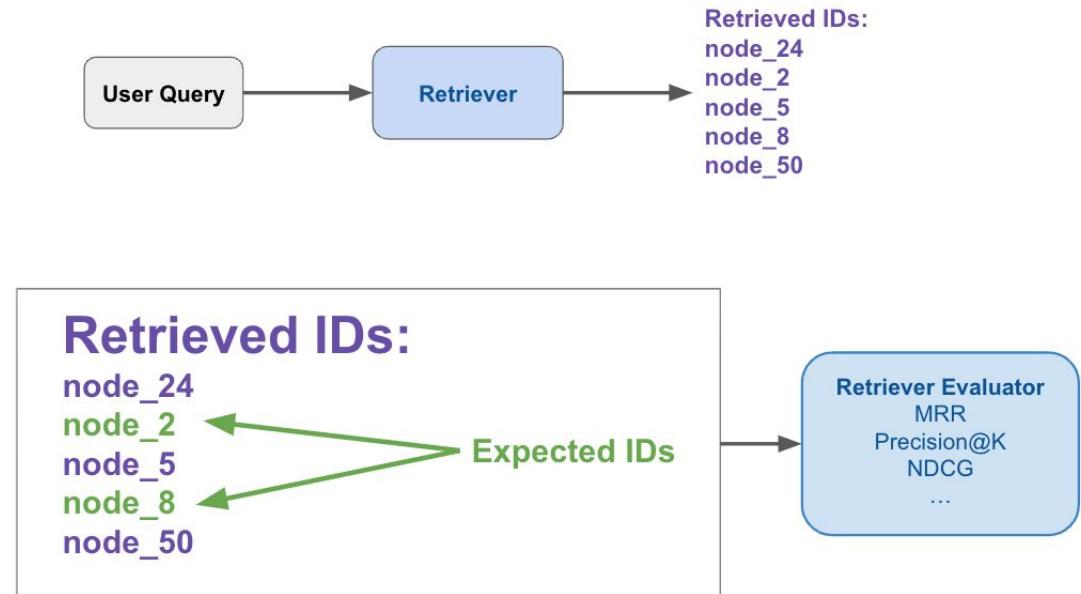
- How do we properly evaluate a RAG system?
 - Evaluate in isolation (retrieval, synthesis)
 - Evaluate e2e
- Open question: which one should we do first?





Evaluation in Isolation (Retrieval)

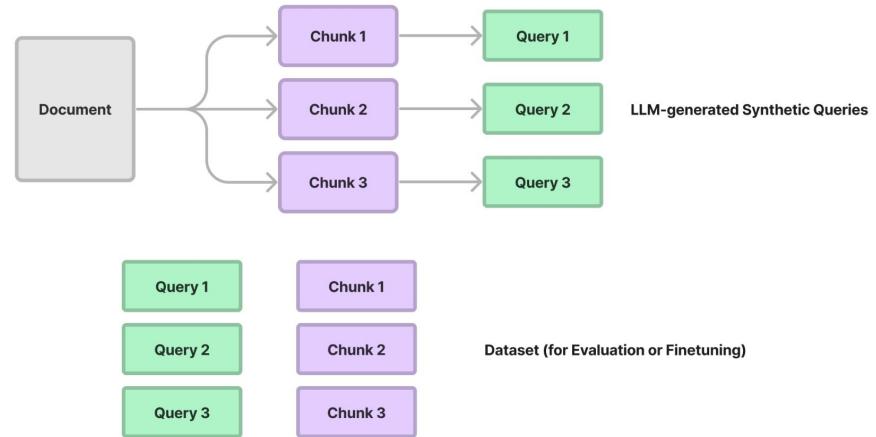
- **Details:** Evaluate quality of retrieved chunks given user query
- **Create dataset**
 - Input: query
 - Output: the “ground-truth” documents relevant to the query
- Run retriever over dataset
- Measure **ranking metrics**
 - Success rate / hit-rate
 - MRR
 - Hit-rate





Synthetic Dataset Generation for Retrieval Evals

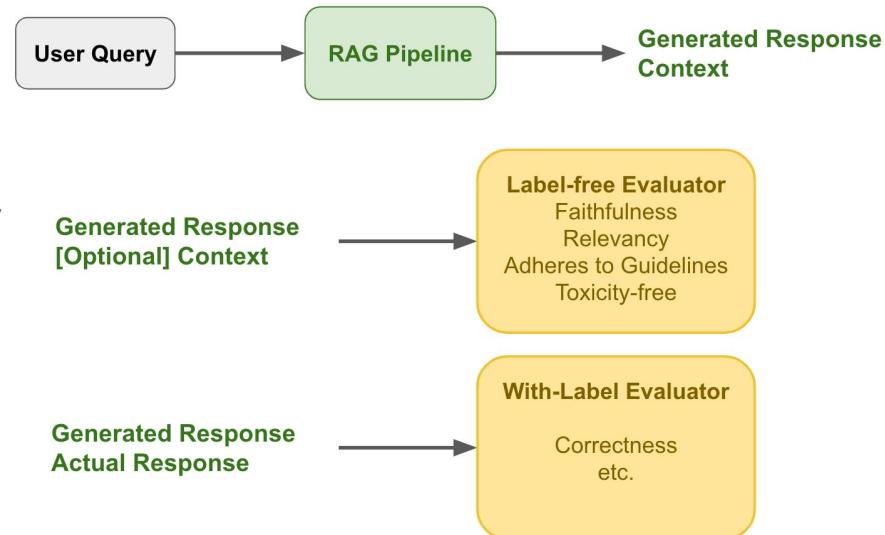
1. Parse / chunk up text corpus
2. Prompt GPT-4 to generate questions from each chunk (or subset of chunks)
3. Each (question, chunk) is now your evaluation pair!





Evaluation E2E

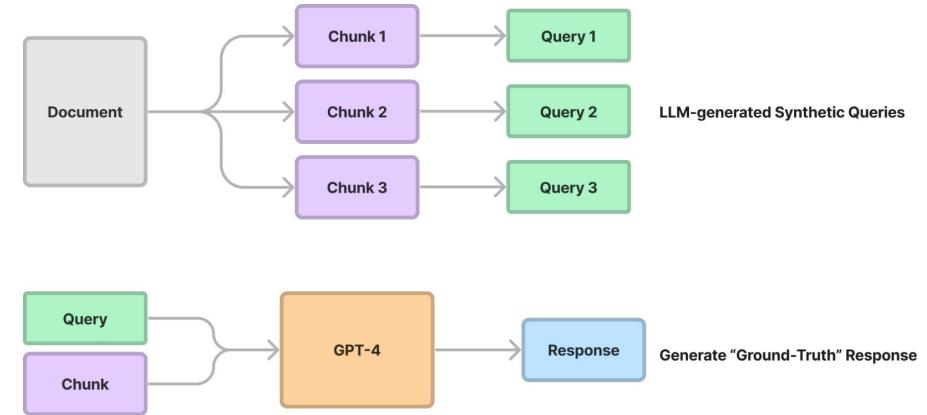
- **Details:** Evaluation of final generated response given input
- **Create Dataset**
 - Input: query
 - [Optional] Output: the “ground-truth” answer
- Run through full RAG pipeline
- Collect evaluation metrics:
 - **If no labels:** label-free evals
 - **If labels:** with-label evals





Synthetic Dataset Generation for E2E Eval

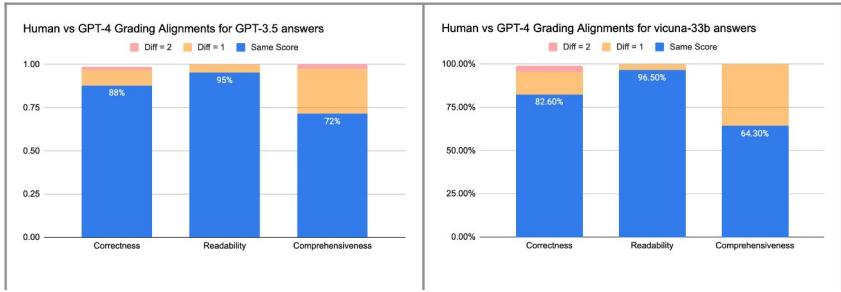
1. Parse / chunk up text corpus
2. Prompt GPT-4 to generate questions from each chunk
3. Run (question, context) through GPT-4
→ Get a “ground-truth” response
4. Each (question, response) is now your evaluation pair!



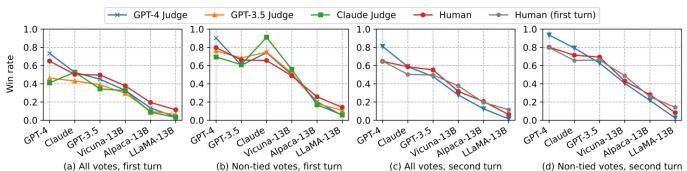


LLM-based Evaluation Modules

- GPT-4 is a good human grader
- **Label-free Modules**
 - **Faithfulness**: whether response matches retrieved context
 - **Relevancy**: whether response matches query
 - **Guidelines**: whether response matches guidelines
- **With-Labels**
 - **Correctness**: whether response matches “golden” answer.



<https://www.databricks.com/blog/LLM-auto-eval-best-practices-RAG>



<https://arxiv.org/pdf/2306.05685.pdf>



Optimizing RAG Systems



From Simple to Advanced

Table Stakes

- Better Parsers
- Chunk Sizes
- Hybrid Search
- Metadata Filters



Less Expressive
Easier to Implement
Lower Latency/Cost

Advanced Retrieval

- Reranking
- Recursive Retrieval
- Embedded Tables
- Small-to-big Retrieval

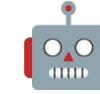


More Expressive
Harder to Implement
Higher Latency/Cost



Fine-tuning

- Embedding fine-tuning
- LLM fine-tuning



Agentic Behavior

- Routing
- Query Planning
- Multi-document Agents

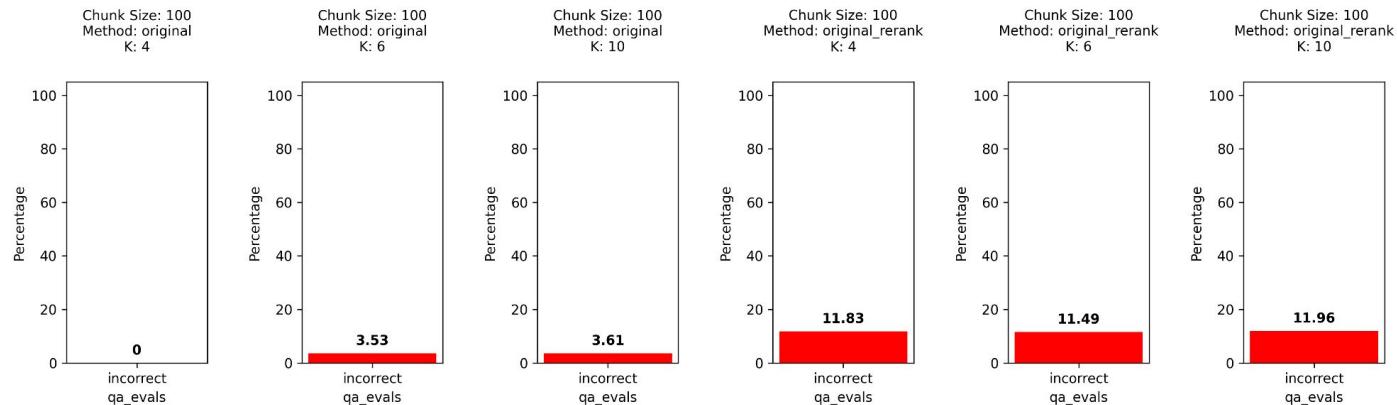


Table Stakes: Chunk Sizes

Tuning your chunk size can have outsized impacts on performance

Not obvious that more retrieved tokens == higher performance!

Note: Reranking (shuffling context order) isn't always beneficial.



Source:
Arize Phoenix + LlamalIndex
Workshop:
<https://colab.research.google.com/drive/1Siuf13rLI-kII1iaNfv-NniBdwUpS?usp=sharing#scrollTo= as7h-u1lwR>



Table Stakes: Prompt Engineering

RAG uses core Question-Answering (QA) prompt templates

Ways you can customize:

- Adding few-shot examples
- Modifying template text
- Adding emotions

Accessing Prompts

Here we get the prompts from the query engine. Note that *all* prompts are returned, including ones used in sub-modules in the query engine. This allows you to centralize a view of these prompts!

```
prompts_dict = query_engine.get_prompts()
```

```
display_prompt_dict(prompts_dict)
```

Prompt Key: response_synthesizer:summary_template

Text:

```
Context information from multiple sources is below.
```

```
-----  
{context_str}  
-----
```

```
Given the information from multiple sources and not prior knowledge, answer the query.
```

```
Query: {query_str}
```

```
Answer:
```



Table Stakes: Customizing LLMs

Task performance on
easy-to-hard tasks
(RAG, agents) varies
wildly among LLMs

Paid LLM APIs

Model Name	Basic Query Engines	Router Query Engine	Sub Question Query Engine	Text2SQL	Pydantic Programs	Data Agents	Notes
gpt-3.5-turbo (openai)	✓	✓	✓	✓	✓	✓	
gpt-3.5-turbo-instruct (openai)	✓	✓	✓	✓	✓	⚠	Tool usage in data-agents seems flakey.
gpt-4 (openai)	✓	✓	✓	✓	✓	✓	
claude-2 (anthropic)	✓	✓	✓	✓	✓	⚠	Prone to hallucinating tool inputs.
claude-instant-1.2 (anthropic)	✓	✓	✓	✓	✓	⚠	Prone to hallucinating tool inputs.

Open Source LLMs

Since open source LLMs require large amounts of resources, the quantization is reported. Quantization is just a method for reducing the size of an LLM by shrinking the accuracy of calculations within the model. Research has shown that up to 4Bit quantization can be achieved for large LLMs without impacting performance too severely.

Model Name	Basic Query Engines	Router Query Engine	SubQuestion Query Engine	Text2SQL	Pydantic Programs	Data Agents	Notes
llama2-chat-7b 4bit (huggingface)	✓	●	●	●	●	⚠	Llama2 seems to be quite chatty, which makes parsing structured outputs difficult. Fine-tuning and prompt engineering likely required for better performance on structured outputs.
Mistral-7B-instruct-v0.1 4bit (huggingface)	✓	●	●	⚠	⚠	⚠	Mistral seems slightly more reliable for structured outputs compared to Llama2. Likely with some prompt engineering, it may do better.
zephyr-7b-alpha (huggingface)	✓	✓	✓	✓	✓	⚠	Overall, zephyr-7b-alpha appears to be more reliable than other open-source models of this size. Although it still hallucinates a bit, especially as an agent.



Table Stakes: Customizing Embeddings

Your embedding model + reranker affects retrieval quality

Embedding	WithoutReranker		bge-reranker-base		bge-reranker-large		Cohere-Reranker	
	Hit Rate	MRR	Hit Rate	MRR	Hit Rate	MRR	Hit Rate	MRR
OpenAI	0.876404	0.718165	0.91573	0.832584	0.910112	0.855805	0.926966	0.86573
bge-large	0.752809	0.597191	0.859551	0.805243	0.865169	0.816011	0.876404	0.822753
llm-embedder	0.814607	0.587266	0.870787	0.80309	0.876404	0.824625	0.882022	0.830243
Cohere-v2	0.780899	0.570506	0.876404	0.798127	0.876404	0.825281	0.876404	0.815543
Cohere-v3	0.825843	0.624532	0.882022	0.806086	0.882022	0.834644	0.88764	0.836049
Voyage	0.831461	0.68736	0.926966	0.837172	0.91573	0.858614	0.91573	0.851217
JinaAI-Small	0.831461	0.614045	0.91573	0.843071	0.926966	0.857303	0.926966	0.868633
JinaAI-Base	0.848315	0.68221	0.938202	0.846348	0.938202	0.868539	0.932584	0.873689
Google-PaLM	0.865169	0.719476	0.910112	0.833708	0.910112	0.85309	0.910112	0.855712

Source:

<https://blog.llamaindex.ai/boosting-rag-picking-the-best-embedding-reranker-models-42d079022e83>



Table Stakes: Metadata Filtering

- **Metadata:** context you can inject into each text chunk
- Examples
 - Page number
 - Document title
 - Summary of adjacent chunks
 - Questions that chunk can answer (reverse HyDE)
- **Benefits**
 - **Can Help Retrieval**
 - **Can Augment Response Quality**
 - **Integrates with Vector DB Metadata Filters**

Example of Metadata

Metadata	Text Chunk
{"page_num": 1, "org": "OpenAI"}	We report the development of GPT-4, a large-scale, multimodal...



Table Stakes: Metadata Filtering

Question: “Can you tell me the risk factors in 2021?”

Raw Semantic Search is **low precision**.

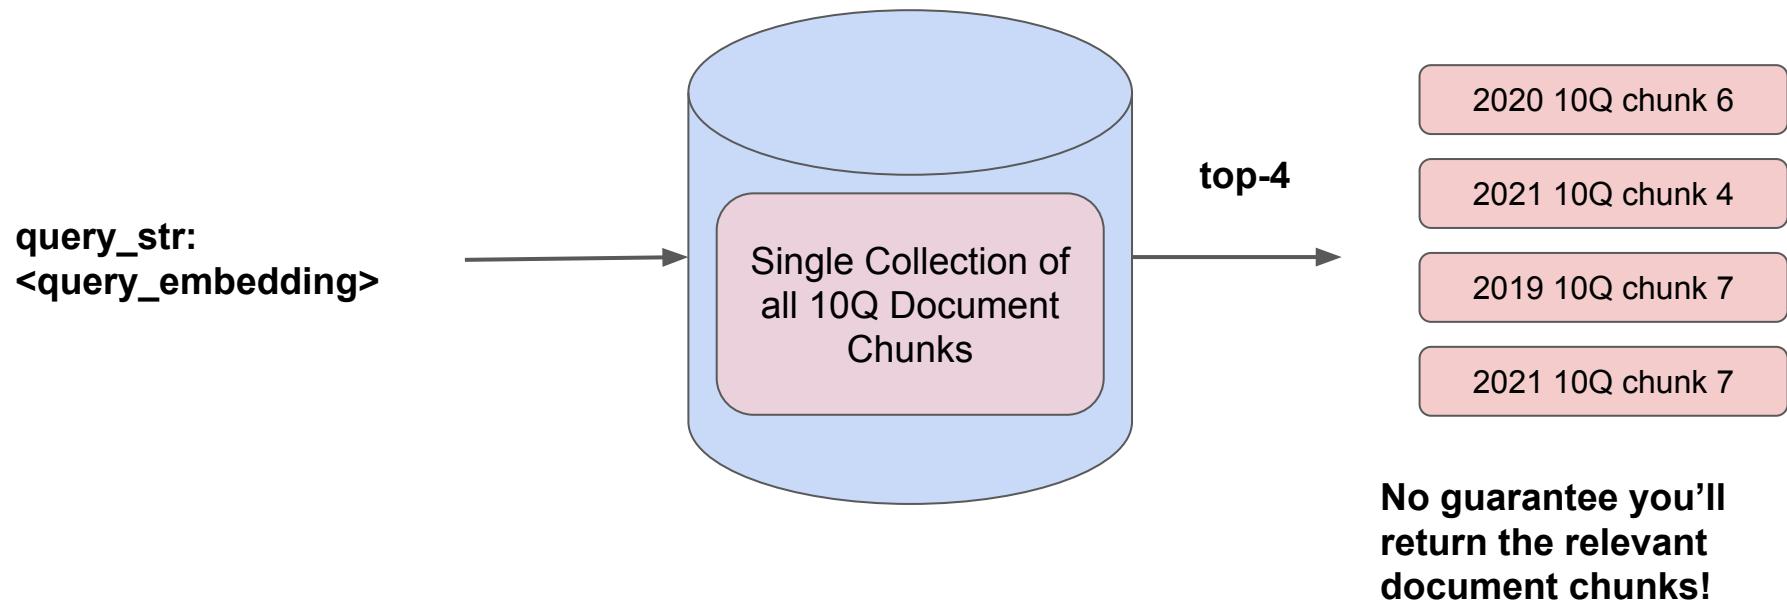




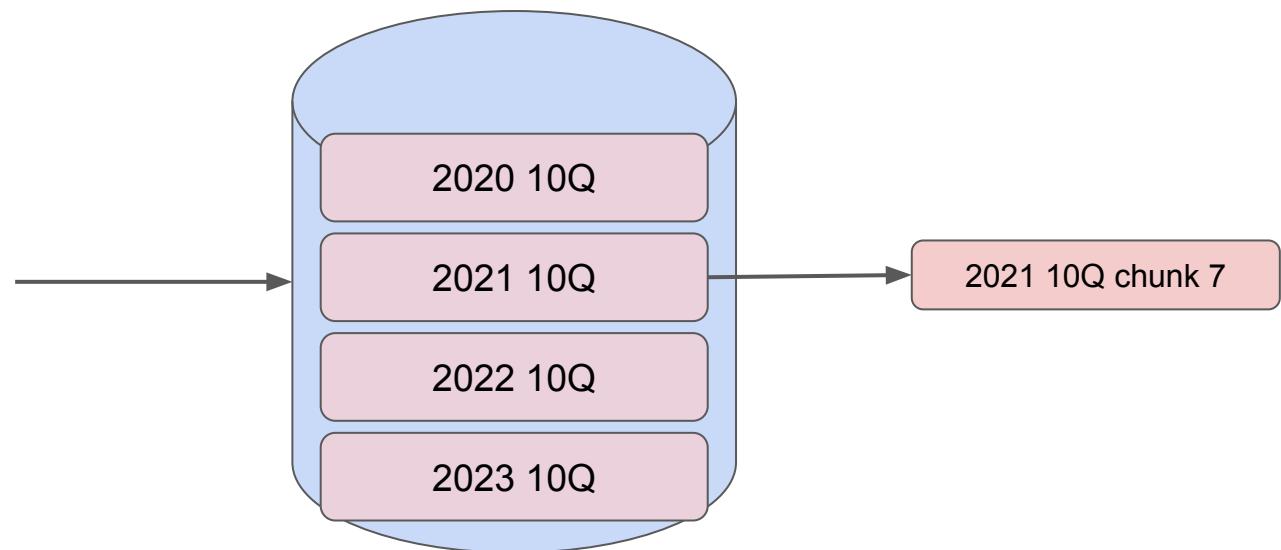
Table Stakes: Metadata Filtering

Question: “Can you tell me the risk factors in 2021?”

If we can *infer* the metadata filters (year=2021), we remove irrelevant candidates, increasing precision!

`query_str:
<query_embedding>`

`Metadata tags:
{“year”: 2021}`





Advanced Retrieval: Small-to-Big

Intuition: Embedding a big text chunk feels suboptimal.

Solution: Embed text at the sentence-level - then **expand** that window during LLM synthesis

Question:
What are the
concerns
surrounding the
AMOC?

Embedding
Lookup

Embed Sentence → Link to Expanded Window

Continuous observation of the Atlantic meridional overturning circulation (AMOC) has improved the understanding of its variability (Frajka-Williams et al., 2019), but there is low confidence in the quantification of AMOC changes in the 20th century because of low agreement in quantitative reconstructed and simulated trends. Direct observational records since the mid-2000s remain too short to determine the relative contributions of internal variability, natural forcing and anthropogenic forcing to AMOC change (high confidence). Over the 21st century, AMOC will very likely decline for all SSP scenarios but will not involve an abrupt collapse before 2100. 3.2.2.4 Sea Ice Changes
Sea ice is a key driver of polar marine life, hosting unique ecosystems and affecting diverse marine organisms and food webs through its impact on light penetration and supplies of nutrients and organic matter (Arrigo, 2014).

What the LLM Sees

What the LLM Sees



Advanced Retrieval: Small-to-Big

Leads to more precise retrieval.

Avoids “lost in the middle” problems.

There is low confidence in the quantification of AMOC changes in the 20th century due to low agreement in quantitative reconstructed and simulated trends. Additionally, direct observational records since the mid-2000s remain too short to determine the relative contributions of internal variability, natural forcing, and anthropogenic forcing to AMOC change. However, it is very likely that AMOC will decline over the 21st century for all SSP scenarios, but there will not be an abrupt collapse before 2100.

Sentence Window Retrieval (k=2)

I'm sorry, but the concerns surrounding the AMOC (Atlantic Meridional Overturning Circulation) are not mentioned in the provided context.

Naive Retrieval (k=5)

Only one out of the 5 chunks is relevant
– “lost in the middle” problem



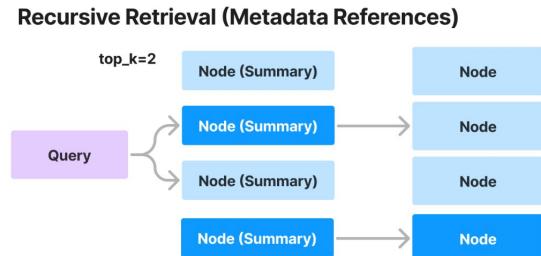
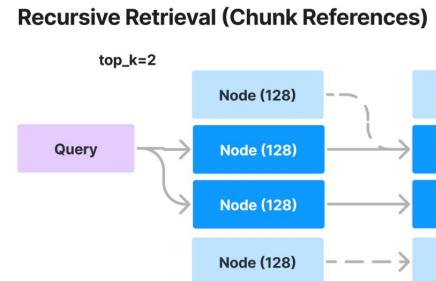
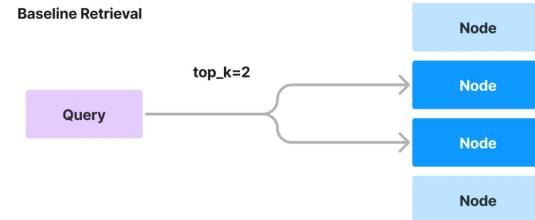
Advanced Retrieval: Small-to-Big

Intuition: Embedding a big text chunk feels suboptimal.

Solution: Embed a smaller **reference** to the parent chunk.
Use parent chunk for synthesis

Examples: Smaller chunks, summaries, metadata

	retrievers	hit_rate	mrr
0	Base Retriever	0.796407	0.605097
1	Retriever (Chunk References)	0.892216	0.739179
2	Retriever (Metadata References)	0.916168	0.746906





Advanced Retrieval Architecture

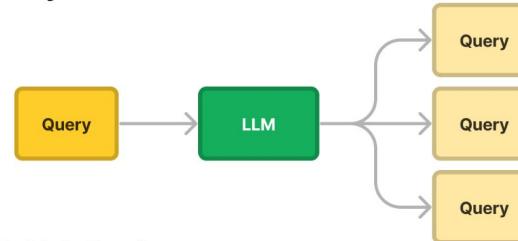


RAG Fusion Retriever

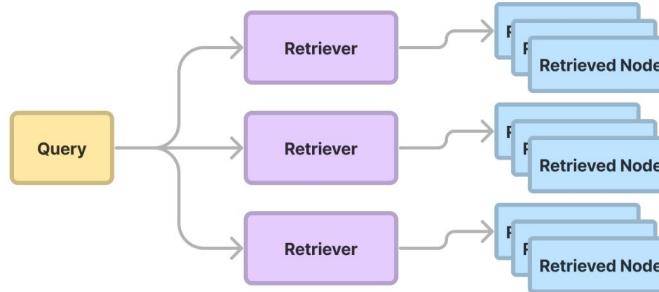
Combines best practices around retrieval:

1. Query Generation/Rewriting
2. Ensemble Retrieval
3. Reranking (with Reciprocal Rank Fusion)
4. Synthesis

1. Query Generation



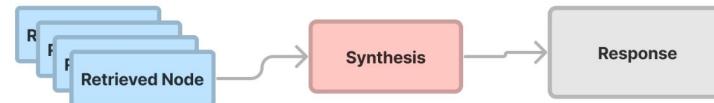
2. Multiple Retrievers



3. Reranking (e.g. Reciprocal Rank Fusion)



4. Synthesis

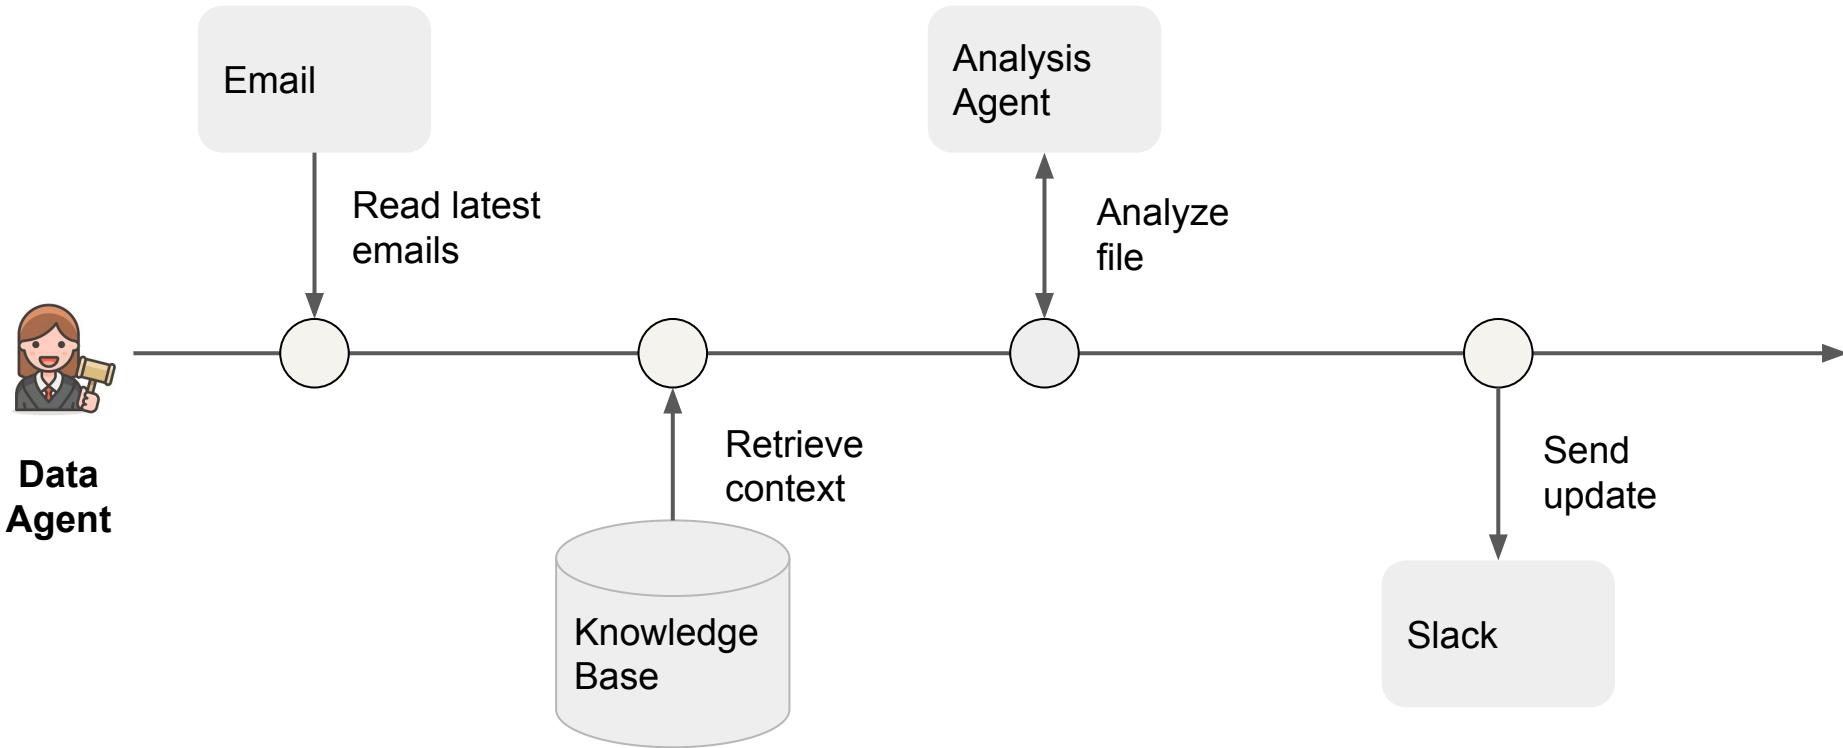




Agents

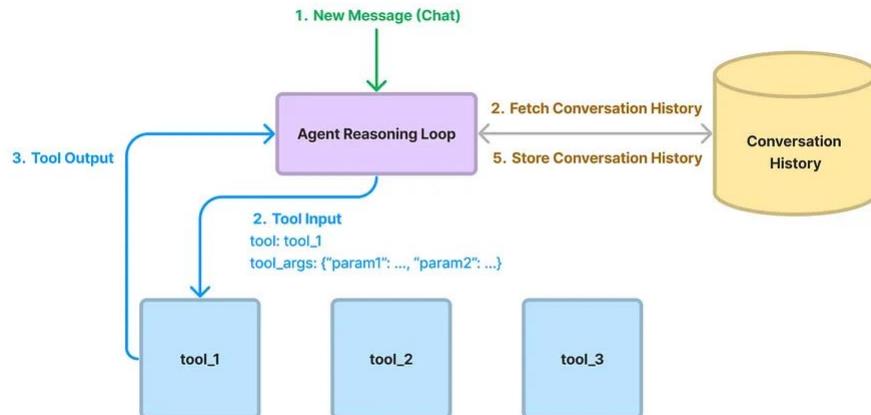


Data Agents - LLM-powered knowledge workers





Data Agents - Core Components



Agent Reasoning Loop

- [ReAct Agent](#) (any LLM)
- [OpenAI Agent](#) (only OAI)

Tools

[Query Engine Tools \(RAG pipeline\)](#)

LlamaHub Tools

- [Code interpreter](#)
- [Slack](#)
- [Notion](#)
- [Zapier](#)
- ... (15+ tools)

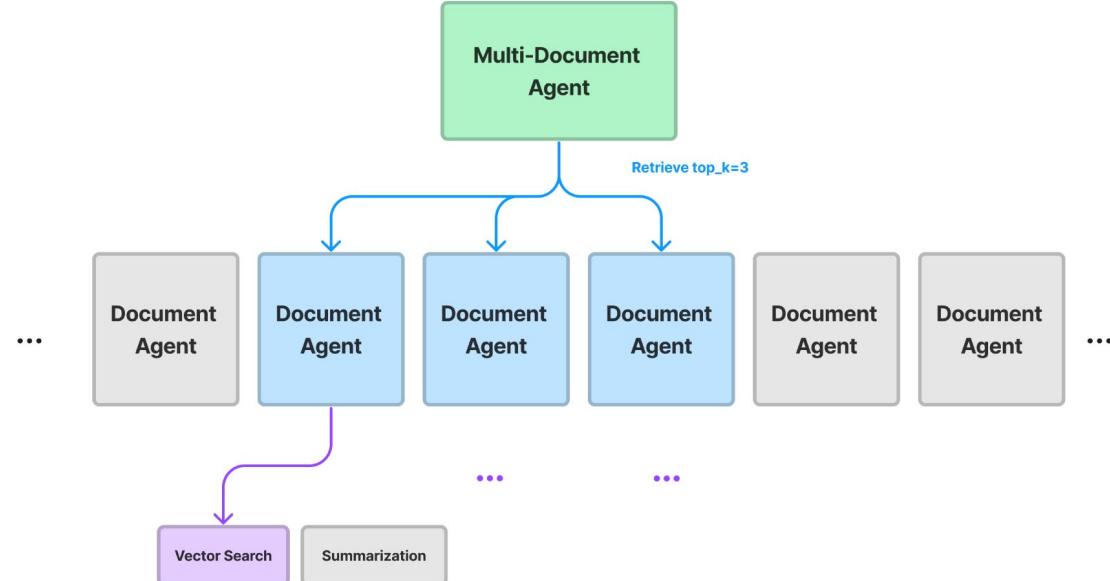


Agentic Behavior: Multi-Document Agents

Intuition: There's certain questions that “top-k” RAG can't answer.

Solution: Multi-Document Agents

- Fact-based QA and Summarization over any subsets of documents
- Chain-of-thought and query planning.





How to retrieve & analyze data from knowledge base?

Use our [query engines](#) as “data tools” over your agent:

- Semantic search
- Summarization
- Text-to-SQL
- Document comparisons
- Combining Structured Data w/ Unstructured

“Simple” Interface - all agent has to infer is a query string!

```
query_engine_tools = [
    QueryEngineTool(
        query_engine=lyft_engine,
        metadata=ToolMetadata(
            name='lyft_10k',
            description="Provides information about Lyft financials for year 2021. "
            "Use a detailed plain text question as input to the tool."
        )
    ),
    QueryEngineTool(
        query_engine=uber_engine,
        metadata=ToolMetadata(
            name='uber_10k',
            description="Provides information about Uber financials for year 2021. "
            "Use a detailed plain text question as input to the tool."
        )
    )
]
```

Example Notebook:

- [OpenAI Agent + query engines \(as tools\)](#)
- [Analyzing structured + unstructured data](#)



Example: Financial Analysis with Agents + RAG

Question: “Compare and contrast Uber and Lyft’s revenue growth”

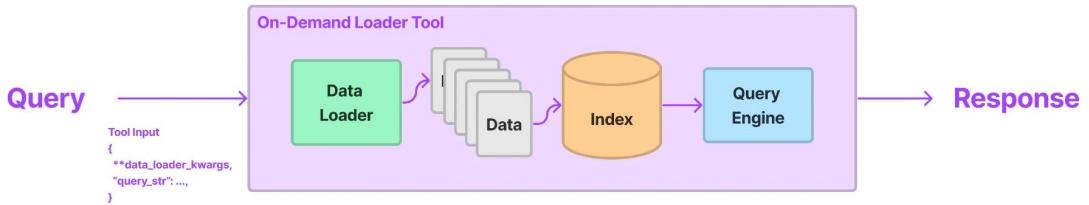
- **Agent:** breaks down question into sub-questions over tools
- **Per-Document RAG:** Answer question over a given document via top-k retrieval.

Let's Try It Out!

```
: agent.chat_repl()  
===== Entering Chat REPL =====  
Type "exit" to exit.  
  
== Calling Function ==  
Calling function: lyft_10k with args: {  
    "input": "What was Lyft's revenue growth in 2021?"  
}  
Got output:  
Lyft's revenue grew by 36% in 2021 compared to the prior year.  
=====  
== Calling Function ==  
Calling function: uber_10k with args: {  
    "input": "What was Uber's revenue growth in 2021?"  
}  
Got output:  
Uber's revenue growth in 2021 was 57%.  
=====  
Assistant: In 2021, Lyft's revenue grew by 36% compared to the previous year, while Uber's revenue growth was higher at 57%. This indicates that Uber experienced a faster rate of revenue growth than Lyft in 2021.
```



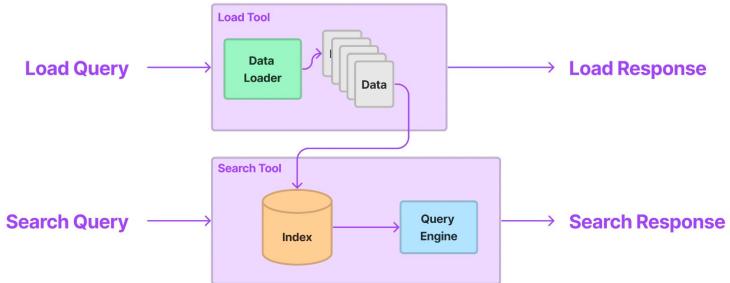
How to handle large responses from tools?



```
from llama_hub.wikipedia.base import WikipediaReader
from llama_index.tools.on_demand_loader_tool import OnDemandLoaderTool

tool = OnDemandLoaderTool.from_defaults(
    reader,
    name="Wikipedia Tool",
    description="A tool for loading data and querying articles from Wikipedia"
)
```

OnDemandLoaderTool



```
from llama_hub.tools.wikipedia.base import WikipediaToolSpec
from llama_index.tools.tool_spec.load_and_search import LoadAndSearchToolSpec

wiki_spec = WikipediaToolSpec()
# Get the search wikipedia tool
tool = wiki_spec.to_tool_list()[1]

# Create the Agent with load/search tools
agent = OpenAIProxy.from_tools(
    LoadAndSearchToolSpec.from_defaults(
        tool
    ).to_tool_list(),
    verbose=True
)
```

LoadAndSearchToolSpec



How to handle large number of tools?

- Build an index over your tools, and retrieve the most relevant ones to pass to your agent.
- [Example Notebook](#)

```
# define an "object" index over these tools
from llama_index import VectorStoreIndex
from llama_index.objects import ObjectIndex, SimpleToolNodeMapping

tool_mapping = SimpleToolNodeMapping.from_objects(all_tools)
obj_index = ObjectIndex.from_objects(
    all_tools,
    tool_mapping,
    VectorStoreIndex,
)
```

```
from llama_index.agent import FnRetrieverOpenAI
```

```
agent = FnRetrieverOpenAI.from_retriever(obj_index.as_retriever(), verbose=True)
```



Fine-Tuning



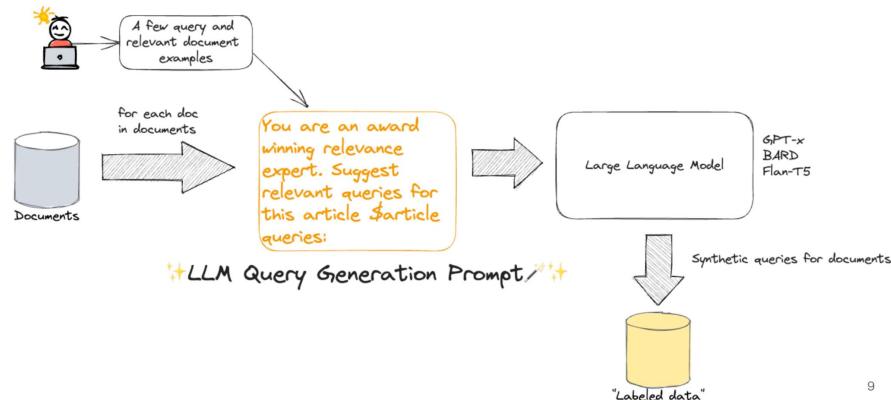
Fine-Tuning: Embeddings

Intuition: Embedding Representations are not optimized over your dataset

Solution: Generate a synthetic query dataset from raw text chunks using LLMs

Use this synthetic dataset to finetune an embedding model.

The gist of using LLMs to generate labeled data

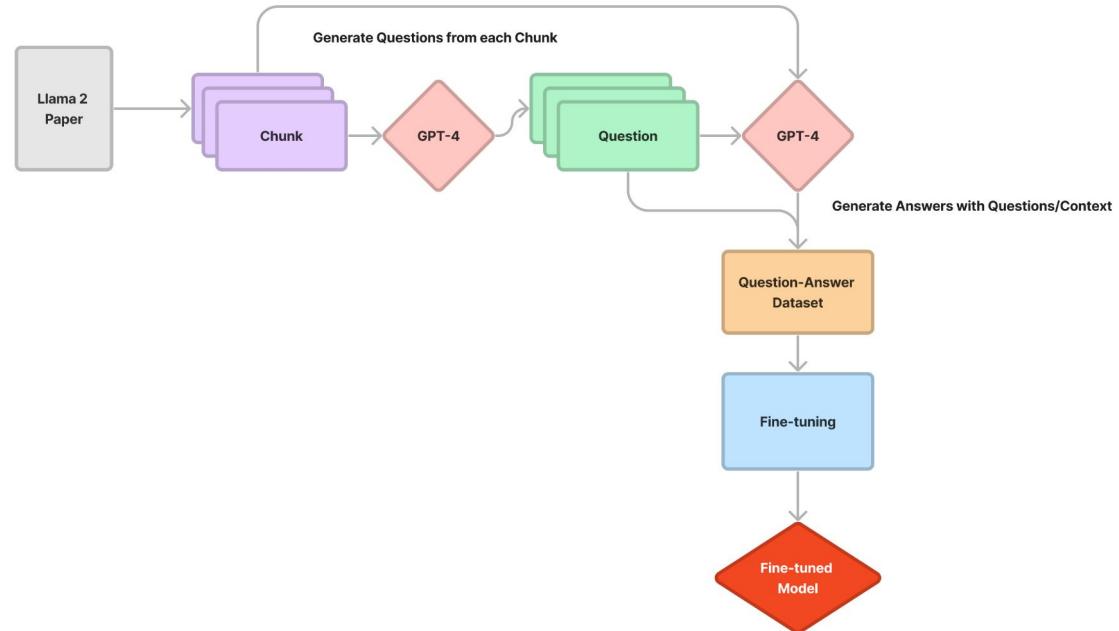




Fine-Tuning: LLMs

Intuition: Weaker LLMs are not bad at response synthesis, reasoning, structured outputs, etc.

Solution: Generate a synthetic dataset from raw chunks (e.g. using GPT-4). Help fix all of the above!





Production RAG Guide

https://gpt-index.readthedocs.io/en/latest/end_to_end_tutorials/dev_practices/production_rag.html





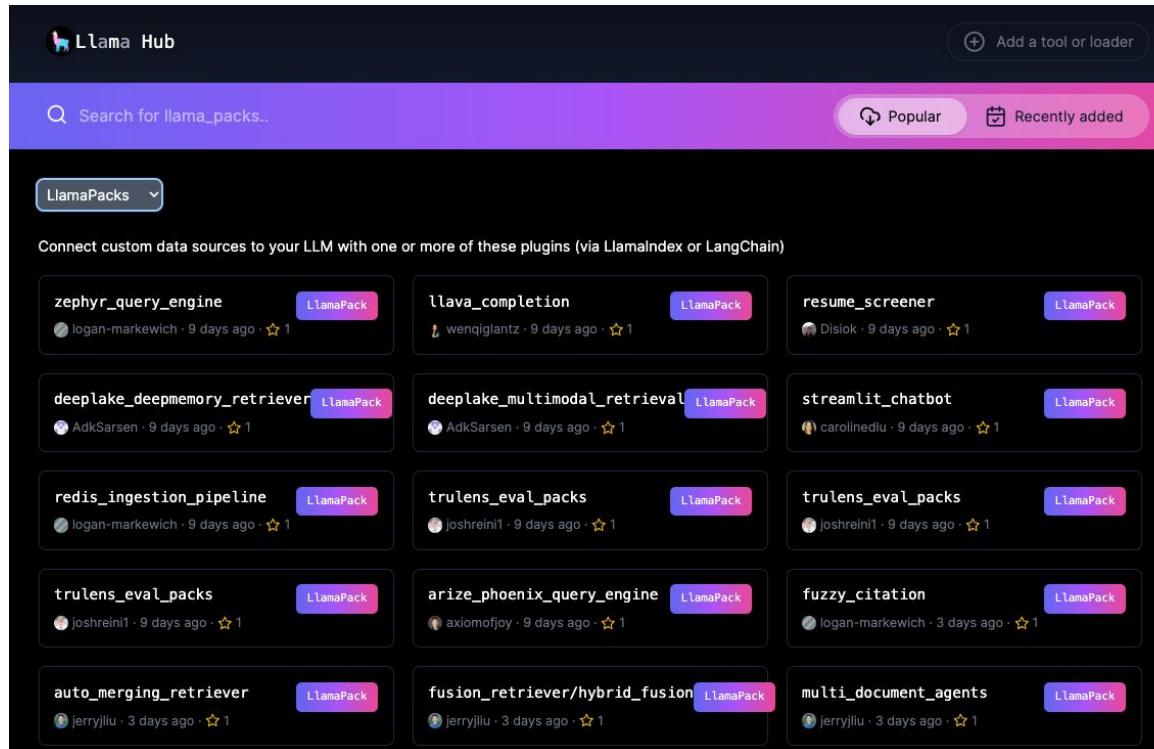
Let's Put it all Together



Kickstart your LLM app

Llama Packs

- A community-driven hub of prepackaged modules
- 25+ diverse packs to get started
- Use it out of the box, or inspect the code



The screenshot shows the Llama Hub interface. At the top, there is a search bar with the placeholder text "Search for llama_packs..", a "Popular" button, and a "Recently added" button. Below the search bar, there is a dropdown menu set to "LlamaPacks". The main area displays a grid of 15 cards, each representing a different LlamaPack. The cards are arranged in five rows and three columns. Each card contains the pack's name, a small icon, the author's name, the date of creation, and a rating of 1 star. The cards are: 1. zephyr_query_engine by logan-markewich (9 days ago) 2. llava_completion by wenqiglantz (9 days ago) 3. resume_screener by Disiok (9 days ago) 4. deeplake_deepmemory_retriever by AdkSarsen (9 days ago) 5. deeplake_multimodal_retrieval by AdkSarsen (9 days ago) 6. streamlit_chatbot by carolinedlu (9 days ago) 7. redis_ingestion_pipeline by logan-markewich (9 days ago) 8. trulens_eval_packs by joshreini1 (9 days ago) 9. trulens_eval_packs by joshreini1 (9 days ago) 10. trulens_eval_packs by joshreini1 (9 days ago) 11. arize_phoenix_query_engine by axiomofjoy (9 days ago) 12. fuzzy_citation by logan-markewich (3 days ago) 13. auto_merging_retriever by jerryliu (3 days ago) 14. fusion_retriever/hybrid_fusion by jerryliu (3 days ago) 15. multi_document_agents by jerryliu (3 days ago)



Build a Full-Stack Application

create-llama

- `create-react-app` but for AI engineers
- Scaffold a full-stack template (choose 3 different backends) with one line of code

Get started by editing `app/page.tsx`

Built by LlamaIndex

what were the risk factors for tesla in 2021?

==== Calling Function ====
Calling function: 2021_tesla_10k with args: {"input":"risk factors"}
Got output: The company is exposed to various risk factors that could have a significant impact on its business, financial condition, and future results. These risks include the potential inadequacy of warranty reserves to cover future warranty claims on products, the possibility that insurance coverage may not be sufficient to protect the company from all business risks, and the presence of covenant restrictions in debt agreements that may limit the company's ability to operate its business. These risk factors underscore the potential challenges and uncertainties that the company may encounter in its operations and industry.
The risk factors for Tesla in 2021 included:
The potential inadequacy of warranty reserves to cover future warranty claims on products.
The possibility that insurance coverage may not be sufficient to protect the company from all business risks.
The presence of covenant restrictions in debt agreements that may limit the company's ability to operate its business.
These risk factors highlight the challenges and uncertainties Tesla may face in its operations and industry.

can you compare and contrast cash flows between 2020 and 2021?

Type a message



SEC Insights

<https://secinsights.ai>

Github:

<https://github.com/run-lla/ma/sec-insights>

← Back to Document Selection Share

06:56 AM What are their main business focus areas?

View progress ▼

Question Received
Generated Sub Query #1

What are the main business focus areas of Amazon?

The main business focus areas of Amazon are serving consumers through online and physical stores, focusing on selection, price, and convenience, offering programs for sellers to grow their businesses, enabling authors, publishers, musicians, filmmakers, and developers to publish and sell content, serving developers and enterprises of all sizes through AWS (Amazon Web Services), and manufacturing and selling electronic devices.

AMZN (2022) p. 3 reflect the way the Company evaluates its business.

AMZN (2022) p. 41 Table of Contents
AMAZON.COM, INC. NOTES T...

AMZN (2022) p. 41 Table of Contents
AMAZON.COM, INC. NOTES T...

Generated Sub Query #2

What are the main business focus areas of Microsoft?

The main business focus areas of Microsoft include reinventing productivity and business processes, building the intelligent cloud and intelligent edge platform, and creating more personal computing.

MSFT (2022) p. 17 continue to attract top talent from across the world. We pla...

MSFT (2022) p. 4 PART I Item 3 Our products include operating systems ...

MSFT (2022) p. 4 PART I Item 3 Our products include operating systems ...

The main business focus areas of Amazon are serving consumers through online and physical stores, offering programs for sellers to grow their businesses, enabling content creators to publish and sell their

AMZN 2022 Table of Contents

AMAZON.COM, INC.

PART I

Item 1. Business

This Annual Report on Form 10-K and the documents incorporated herein by reference contain forward-looking statements based on existing information and current expectations. Actual results and outcomes may differ materially from those expressed in forward-looking statements. See Item 1A of Part I — “Risk Factors.” As used herein, “Amazon.com,” “we,” “our,” and similar terms include Amazon.com, Inc. and its subsidiaries, unless the context indicates otherwise.

General

We seek to be Earth’s most customer-centric company. We are guided by four principles: customer obsession rather than competitor focus, passion for invention, commitment to operational excellence, and long-term thinking. In each of our segments, we serve our primary customer sets, consisting of consumers, sellers, developers, enterprises, content creators, advertisers, and employees.

We have organized our operations into three segments: North America, International, and Amazon Web Services (“AWS”). These segments reflect the way the Company evaluates its business performance and manages its operations. Information on our net sales is contained in Item 8 of Part II, “Financial Statements and Supplementary Data — Note 10 — Segment Information.”

Consumers

We serve consumers through our online and physical stores and focus on selection, price, and convenience. We design our stores to enable hundreds of millions of unique products to be sold by us and by third parties across dozens of product categories. Customers access our offerings through our websites, mobile apps, Alexa, devices, streaming, and physically visiting our stores. We also manufacture and sell electronic devices, including Kindle, Fire tablet, Fire TV, Echo, Ring, Bling, andero, and we develop and produce media content. We seek to offer our customers low prices, fast and free delivery, easy returns functionality, and timely customer service. In addition, we offer subscription services such as Amazon Prime membership program that includes fast, free shipping on millions of items, access to award-winning movie and music services, and other benefits.

We fulfill customer orders in a number of ways, including through: North America and International fulfillment networks that we operate; co-sourced and outsourced arrangements in certain countries; digital delivery; and through our physical stores. We operate customer service centers globally, which are supplemented by co-sourced arrangements. See Item 2 of Part I, “Properties.”

Sellers

We offer programs that enable sellers to grow their businesses, sell their products in our stores, and fulfill orders through us. We are not the seller of record in these transactions. We earn fixed fees, a percentage of sales, per-unit activity fees, interest, or some combination thereof, for our seller programs.

Developers and Enterprises

We serve developers and enterprises of all sizes, including start-ups, government agencies, and academic institutions, through AWS, which offers a broad set of on-demand technology services, including compute, storage, database, analytics, and machine learning, and other services.

Content Creators

We offer programs that allow authors, independent publishers, musicians, filmmakers, Twitch streamers, skill and app developers, and others to publish and sell their content.

Advertisers

We provide advertising services to sellers, vendors, publishers, authors, and others, through programs such as sponsored ads, display, and video advertising.

3



Thanks!

Building LLMs in an enterprise setting? We'd love to chat!

